How does a chatbot actually work, and what does it take to build one?

This book follows a large language model from start to finish, in plain English with no math: the data it reads, the chips and data centers it's trained on, the electricity and water it uses, and the training that turns a word predictor into a helpful assistant.

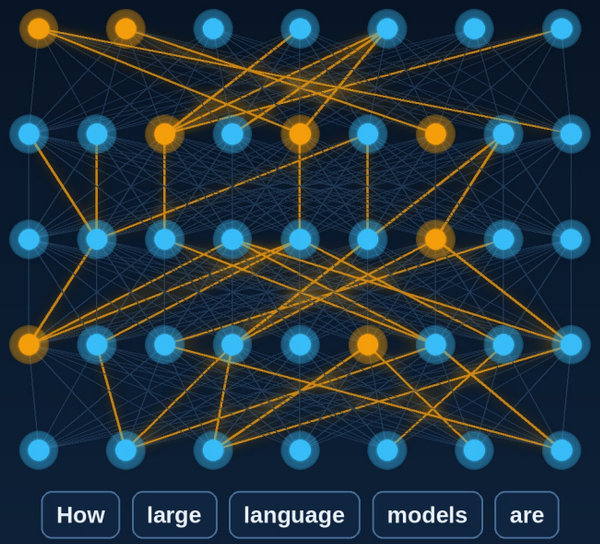
- Tokens, embeddings and the Transformer, explained with everyday analogies
- GPUs, TPUs, AI factories, and the real numbers on energy and water
- Post-training, reasoning models, benchmarks, validation and safety
- A guided tour of the open models (Llama, Qwen, DeepSeek, Mistral, Gemma and more) and the frontier labs
- Models that see, draw, make video, speak and act as agents

Dozens of surprising facts, every one sourced, with credit to the people who invented each idea. Current as of October 2026.

Aadhavan Muthurengan is an automation and AI professional with more than 15 years of experience in the automotive industry, where Aadhavan has built in-house tools that transformed product development and validation.

HOW LARGE LANGUAGE MODELS ARE MADE

AADHAVAN MUTHURENGAN



HOW LARGE LANGUAGE MODELS ARE MADE

THE COMPLETE STORY, END TO END

From Data and Chips to Chatbots, Reasoning, Images, Video and Agents, in Plain English

NO MATH NEEDED • 27 CLEAR DIAGRAMS • EVERY FACT SOURCED

AADHAVAN MUTHURENGAN